

Queensland University of Technology
Faculty of Science
School of Mathematical Sciences

PhD Stage 2 Proposal

Computational Modelling of Multilateral Treaties

Candidate: Oriolson Rodriguez Ramirez
Student Number: n12608521
Degree: Doctor of Philosophy (PhD)
Mode of Study: Full-time
Submission Date: July 2, 2026

Submitted in partial fulfilment of the requirements
of the degree of Doctor of Philosophy at
Queensland University of Technology.

Contents

proposed title	1
Supervisory Team	1
1 Background and Literature Review	2
1.1 Introduction	2
1.2 The Antarctic Treaty System as a consensus regime	3
1.3 Quantitative and text-based study of institutions and treaties	4
1.4 Measuring policy preference from political text	5
1.5 Modelling consensus: responsive and anticipatory approaches	7
1.6 Negotiation and coalition formation simulation using reinforcement learning	8
1.7 Summary of literature gap	9
1.8 Research problem	10
2 Program and design of the research project	11
2.1 Objectives, methodology and research plan	11
2.1.1 Objective 1 / Chapter 1: The Support layer	11
2.1.2 Objective 2 / Chapter 2: The Opposition layer and external validation	15
2.1.3 Objective 3 / Chapter 3: The Simulation layer	17
2.2 Resources and funding required	19
2.3 Individual contribution to the research team	20
2.4 Acknowledging use of Artificial Intelligence (AI) and AI-assisted tech- nologies	20

2.5 Timeline for completion 20

References 24

Appendices 28

Appendix A: Research Integrity Online 28

Appendix B: Coursework 28

Appendix C: Notation 28

List of Figures

2.1	Thesis workflow: from text records to a validated consensus simulation.	12
2.2	Chapter 1 support-layer pipeline, from text records to the country $\times$ topic $\times$ year support tensor S	16
2.3	Chapter 2 opposition layer sketch: from text data to opposition tensor O , forming the signed preference network (S, O)	18

Proposed thesis title and type

“Computational Modelling of Multilateral Treaties”, thesis by publication.

Proposed supervisory team and their credentials

Role	Name	Affiliation
Principal Supervisor	Prof. Michael Bode	School of Mathematical Sciences, QUT
Associate Supervisor	Prof. Jonathan Rhodes	School of Mathematical Sciences, QUT

Principal Supervisor — Prof. Michael Bode Prof. Bode is Deputy Director of Capability Development at SAEF (Securing Antarctica’s Environmental Future). His work centres on applied mathematics and conservation, with a focus on decision-making under uncertainty and mathematical biology. His previous work on Antarctic Treaty System (ATS) decision-making and conservation makes him an expert on consensus-based systems.

Associate Supervisor — Prof. Jonathan Rhodes Prof. Rhodes is Director of the Centre for Environment and Society. He specialises in spatial ecology, environmental governance, and the modelling of coupled human–natural systems, bridging quantitative models with real-world policy. His expertise gives him deep insight into the policy implications of environmental decisions, providing a grounding for the treaty-system applications explored in this research.

Chapter 1

Background and Literature Review

1.1 Introduction

Consensus is a collective decision-making process where every party can object to any proposal and effectively veto it. It becomes a consensus regime when consensus becomes the only means to reach decisions. Consensus is not the same as unanimity: a unanimity rule puts a proposal to a vote and records who votes in favour, against or abstain, whereas under consensus a decision proceeds whenever no party registers a formal objection (Rietig et al. 2023). Since in such regimes a single party can halt any proposal even when no other party objects, an outside observer might, at first sight, expect such regimes to decide very little. This is even more pronounced when there is a large and heterogeneous group: the larger the group, the smaller the set of proposals that change the status quo (Tsebelis 2002). In practice, consensus regimes can reach agreement even on contested, transnational issues, often by relying on informal practices: trust, dialogue, and leadership. These practices fill the gaps left by their formal rules (Sykora-Bodie and Morrison 2019). But success is never guaranteed: in bodies such as the United Nations Security Council a single veto routinely blocks collective action, and even historically productive regimes can stall (Gardiner et al. 2025). Quantitatively, the picture is no different: what separates a consensus regime that keeps reaching decisions from one that can no longer reach them is not captured by existing models of consensus dynamics (Pérez et al. 2018; Tran et al. 2024). A model that better reproduces how consensus is reached or broken is therefore needed, both to explain past behaviour and to anticipate when a regime approaches the limits of its capacity to act.

Models of multi-actor decision-making already exist, but they were not built for this task. Classical opinion-dynamics and game-theoretic models represent each actor with a single, low-dimensional position and are rarely use the actual text those actors produce, such as their statements and negotiating documents. Capturing a consensus regime

faithfully requires more: a higher-dimensional description in which actors hold distinct positions across many issues, preferences read directly from the text the actors produce and organised by topic, and measures validated and calibrated against observed outputs rather than assumed. Recent advances in natural language processing and machine learning make each of these upgrades feasible.

The Antarctic Treaty System (ATS) is one of those consensus-driven regimes. At its centre is the Antarctic Treaty (AT), whose decision-making parties are called consultative parties (Secretariat of the Antarctic Treaty 2023). These parties meet yearly at the Antarctic Treaty Consultative Meeting (ATCM), where outputs are adopted only by consensus. As international relations strain and Antarctic environmental pressures intensify, the risk of losing the capacity to act increases. In 2022, for the first time, the ATCM parties failed to reach consensus on the final report; even as activity has increased, the regime has become less responsive to challenges (Gardiner et al. 2025).

1.2 The Antarctic Treaty System as a consensus regime

At the ATCM, Consultative Parties (CPs) reach decisions exclusively by consensus; non-consultative parties (NCPs), observers, and experts attend but take no part. These outputs take three forms (Secretariat of the Antarctic Treaty 2023, Rule 24). **Measures** set out substantive provisions that the parties intend to be legally binding in their domestic law; they require each party's government to formally approve them. **Decisions** settle the ATCM's own internal and organisational matters. **Resolutions** are non-binding: they record recommendations or expressions of intent, with no legal force.

There are four types of input:

- **Working Papers (WP)** (Secretariat of the Antarctic Treaty 2023, Rule 48): submitted by CPs and Observers, these require discussion at the ATCM and are the main source of outputs. They are the only input from CPs that can produce an outcome.
- **Secretariat Papers** (Secretariat of the Antarctic Treaty 2023, Rule 49): organisational matters such as the budget.
- **Information Papers** (Secretariat of the Antarctic Treaty 2023, Rule 50): supporting information, also from NCPs and experts.

- **Background Papers** (Secretariat of the Antarctic Treaty 2023, Rule 51): documents submitted for the record only.

The consensus system makes the ATS a notable example of international cooperation (Jackson 2021); however, the same mechanism has become a major challenge as the number of CPs grows (Tsebelis 2002; Gardiner et al. 2025).

A notable case of contested consensus occurred in the late 1970s and 1980s, during the negotiations over mineral resources in Antarctica. The Convention on the Regulation of Antarctic Mineral Resource Activities (CRAMRA) was a framework, agreed in 1988, that would have permitted and regulated mining in Antarctica. It was abandoned before entering into force and replaced by the Madrid Protocol (1991), which instead prohibited all mineral resource extraction and designated Antarctica a natural reserve devoted to peace and science (Jackson 2021). Although it unfolded in two stages, it forms a single documented episode: one negotiation over Antarctic mineral resources in which the parties' positions are recorded and then reversed. That documented reversal makes it a valuable canary test, a known case against which any quantitative stance measure can be validated.

1.3 Quantitative and text-based study of institutions and treaties

The quantitative study of international institutions long predates its application to Antarctica. Substantial literature codes the design of treaties and organisations as structured variables: categorical or ordinal attributes such as their decision rules, dispute-settlement provisions, and enforcement mechanisms. Some studies link these variables to institutional performance, survival, and decline (Gray 2018). A more recent trend treats the *text* of international agreements and debates as data (Grimmer and Stewart 2013). They recover state positions from voting records and speeches (Baturo et al. 2017; Slapin and Proksch 2008). Together, these approaches show that the preferences and design features of treaty regimes can be measured usefully, not merely described: such measures track institutional outcomes and recover known positions.

The ATS's health has been assessed qualitatively (Mancilla and Jabour 2022) and quantitatively. Quantitatively, it has been analysed in two ways. Gardiner et al. (2025) provides performance accounts that measure the regime in aggregate, counting inputs and outputs, and the time legally binding measures take to enter into force. It concludes that responsiveness to challenges in the Antarctic region is declining. A Working Paper may be submitted by a single party or co-sponsored by several. The second

methodology is structural: it uses this co-sponsorship to map who collaborates with whom, through co-authorship networks of ATCM documents (Botelho et al. 2026), and how such networks form (Kumar et al. 2026). This thesis uses both lines of work. The stances it reads from the text should explain the co-sponsorship behaviour these networks describe: parties with similar stances are likely to co-sponsor papers together, and Section 2.1.2 turns this into a validation check against observed co-sponsorship. The same measures should also help predict the regime’s output series reported by Gardiner et al. (2025), and Chapter 2 tests this. Neither analysis, however, considers what the papers are *about*: both count documents or map their sponsors without reading the content. As a result, they cannot say which issues divide the parties, a limitation taken up in Section 1.7.

1.4 Measuring policy preference from political text

Positional measures place each actor on a policy scale, for example from favouring regulation to favouring access, using only the text the actor produces. Extracting them from political texts is a well-developed field with two paradigms. They differ in how they turn text into numbers: the first paradigm uses word counts (i.e. frequencies), while the second uses embeddings. The first paradigm is scaling and ideal-point methods, which recover a latent positional dimension from the data itself: the axis is not chosen in advance, it emerges from patterns in the text, and its meaning (e.g. regulation versus access) is interpreted after estimation. Wordfish (Slapin and Proksch 2008) is a Poisson scaling model of party text. Its aim was to estimate party positions over time without hand-coded reference texts. The model treats each word count as a Poisson draw whose expected frequency depends on the party’s position on a single latent dimension, and it recovered credible left–right trajectories for German parties from their election manifestos (1990–2005). Its limitation is that the recovered dimension is uninterpreted: the model finds the dominant axis of variation in word usage, which need not correspond to any policy dimension of interest. In a corpus as topically diverse as the ATCM Working Papers, that axis would likely reflect what parties talk about rather than what they support. Lauderdale and Herzog (2016) addressed the problem that legislators talk about different things in different debates. Their model scales each debate separately with Wordfish and then combines the debate-level positions into a general dimension through a Bayesian factor analysis. Applied to speeches in the Irish Dáil and the US Senate, it recovered within-party disagreement that voting records hide. The approach still inherits the uninterpreted-dimension problem: each debate scale is whatever axis dominates that debate’s word frequencies, which need not be the substantive for-or-against axis of interest. Lauderdale and Herzog get away with this because legislative

debates are adversarial, so the dominant axis is usually government versus opposition; ATCM Working Papers have no such structure. Two things are therefore missing for this thesis: the topics have to be discovered from the papers themselves rather than given in advance, and the output must be a stance per topic rather than one general dimension. Vafa et al. (2020) produced a probabilistic framework that jointly places ideal points and models actors' language. The closest dataset to the ATCM corpus is the one Baturo et al. (2017) use to derive state-level preferences from United Nations (UN) debate speeches. Their corpus collects over 7300 annual General Debate statements (1970–2014), one per country per year, in a setting that parallels the ATCM: formal, annual, state-level, and translated into official languages. They derive preferences in two ways, both from word frequencies: Wordscores, a supervised scaling in which the analyst defines the axis in advance through reference texts (US and Russian statements anchor a US–Russia dimension), and correspondence analysis, an unsupervised reduction whose dimensions are interpreted afterwards. Finally, Parschan and Jakob (2025) review the text-based ideal-point estimation algorithms of this first paradigm, comparing their assumptions and variants in one place. The second paradigm uses embeddings as its main tool. An embedding maps a piece of text (a word, a sentence, or a whole document) to a vector in $\mathbb{R}^d$, with d typically in the hundreds. The map is learned from large corpora so that texts appearing in similar contexts receive nearby vectors; comparisons of meaning then become geometric operations, such as distances and angles between vectors. Modern embeddings are produced by neural language models such as BERT (Devlin et al. 2019), which make the vectors context-dependent: the same word receives different vectors in different sentences. Rheault and Cochrane (2019) show that word-embedding geometry recovers ideological placement using parliamentary data. Rodriguez and Spirling (2022) provide a comprehensive analysis of what works and what does not when analysing political text using embedding techniques. They find embeddings useful for analysing political text, but caution that the choice of embeddings strongly affects applied inference, and that the results must be validated.

This thesis adopts the second paradigm and uses embeddings at several stages: to discover what topics the Working Papers discuss, to measure how relevant each paper is to each topic, and, through the same family of neural models, to read each paper's stance on a topic. It also differs from the ideal-point methods in an important way. Ideal-point models summarise each actor with one number: a single position on one latent axis, assumed stable across issues. That summary is too coarse for the ATCM, where a party can, for example, support stricter environmental protection yet oppose new tourism rules. This project therefore requires a topic-conditioned stance: for each party, a separate position on each topic, allowed to change through time. The topic structure is discovered using BERTopic (Grootendorst 2022), which clusters document embeddings and labels each cluster with its most distinctive words. The stance is

read with natural language inference (NLI) models (Burnham 2023; Abercrombie and Batista-Navarro 2022): classifiers that take a text and a stated hypothesis (e.g. “this paper supports stricter protection”) and judge whether the text entails, contradicts, or is neutral towards it. The results can be stress-tested in three ways. The discovered topics are compared against the categories that the Antarctic Treaty Secretariat assigns to each WP by hand. The stances are validated against the CRAMRA canary test: applied to the papers of the minerals negotiations, the pipeline should recover the documented positions, including the late reversal by Australia and France (Jackson 2021). Both can also be tested against observed co-sponsorship. This thesis assumes that countries co-sponsor a WP because they share its topic and stance. That assumption is testable: within the same year and topic, countries whose recovered stances are similar should be the ones that actually co-sponsor papers, with stance similarity computed excluding the paper being predicted. Section 2.1.2 sets out these validation checks.

1.5 Modelling consensus: responsive and anticipatory approaches

There are two traditional groups of models for how groups reach consensus. The first group is responsive and interactive: each actor repeatedly reacts to the current positions of the others. These are the synchronization and opinion-dynamics models; examples are the Kuramoto and opinion-changing-rate models used by Botelho et al. (2026), which represent agreement as a coupled-oscillator convergence. These methods treat consensus as *convergence*. DeGroot (1974) iteratively average agents’ opinions using $x(t+1) = Wx(t)$ until they converge. A bounded-confidence version was later developed by Hegselmann and Krause (2002). A further model in this group is the stubborn-agents model (Friedkin and Johnsen 1990), in which each agent only partially updates toward its neighbours and stays anchored to its initial opinion. Opinion-dynamics models can nevertheless test for agreement: declare consensus when all final positions lie within a tolerance radius, as Bernardo et al. (2021) do for the negotiation of the 2015 Paris Agreement. That test, however, captures preference convergence, not the consensus rule the ATS uses: a procedural right under which one party can block a decision no matter how close the remaining positions are. These opinion-dynamics models also contain no strategic preference, no veto, and no learning; the bargaining behaviour of consensus regimes is therefore outside their scope.

The second group is anticipatory: each actor reasons about how the others will respond and chooses its actions to maximise an objective function. Responsive agents, by contrast, follow an update rule with no objective at all. Anticipatory models use social

choice and game theory to treat consensus as a unanimity rule. Tsebelis (2002) develop a veto-players framework in which every party has veto power. Let q denote the status quo, p a proposal, and u_i the utility function of party i , which scores how much that party values each policy. The status quo is replaced only when $u_i(p) \geq u_i(q)$ for every party i : no party may be left worse off than it is now. Consequently, as the number of parties increases and their positions spread, the set of adoptable policies shrinks and the system becomes harder to change. Analysing the ATS, Gardiner et al. (2025) point to the same ingredients: a growing and diversifying membership, increasingly divergent views, and a consensus rule, and view the regime’s shift away from binding agreements as symptomatic of declining decision-making performance.

Each group has strengths and limits. Responsive models are simple, can be scaled to many agents, and are easy to calibrate against data. But their agents have no objective, no veto, and no strategy; therefore, they capture convergence rather than bargaining. Anticipatory models represent vetoes and strategic choice directly, and they explain *why* a proposal passes or fails. But their equilibria become hard to compute as parties and issues multiply, and they are rarely confronted with data from real negotiations.

Between these two groups lies a middle ground: models that retain the strategic, veto-bearing structure of the second group but, instead of solving for equilibrium in closed form, learn it from repeated interaction. This middle ground has grown into a field of its own and is treated separately in the next section.

1.6 Negotiation and coalition formation simulation using reinforcement learning

Section 1.5 left two requirements for a model of ATS behaviour. The model must represent many heterogeneous parties, each holding a veto and bargaining over a status quo. It must also remain computable at that scale, which the closed-form equilibria of analytic game theory are not. Multi-agent reinforcement learning (MARL) meets both requirements: it keeps the game-theoretic substance but learns equilibrium behaviour from repeated interaction instead of solving for it.

Formally, MARL is defined over stochastic (Markov) games, a generalisation of repeated games in which the state of the negotiation changes as the agents act. Each agent observes the state, chooses an action, and receives a reward: the objective function it learns to maximise. The learning procedures approximate equilibrium behaviour over many simulated rounds rather than deriving it in closed form (Littman 1994). Neural networks (NN) enter only as function approximators, allowing this learning to scale to large state and action spaces.

Foundational MARL work covers multi-issue bargaining (Chang 2021), negotiating-team formation (Bachrach et al. 2020), bargaining under risk and asymmetric information (Schmid et al. 2020), and search-improved game-theoretic learning (Li et al. 2023). Two reviews cover the field: Gronauer and Diepold (2022) on MARL and Sarkar et al. (2022) on coalition formation. Most recently, MARL agents based on Large Language Models (LLMs) have been used to model political negotiation directly: for example, Moghimifar et al. (2024) simulate coalition formation in parliament using text data from party manifestos. Bolleddu (2025) present a MARL system for consensus building and negotiation based on synthetic data. Kulkarni et al. (2025) use negotiated agreements with game-theoretic reasoning to predict dynamic coalition formation in Diplomacy, a board game in which seven powers negotiate, form, and break alliances, widely used as a benchmark for this type of simulation. Finally, some other benchmarks for real negotiation data are starting to appear (Benac et al. 2026). However, these simulations run on synthetic data or abstract benchmarks rather than the deliberative records of a real consensus regime.

1.7 Summary of literature gap

The descriptive work, the aggregate performance accounts and qualitative health assessments of Section 1.3 (Gardiner et al. 2025; Mancilla and Jabour 2022), shows what the regime outputs but not who produces what. The structural analyses of co-authorship networks (Botelho et al. 2026; Kumar et al. 2026) are undirected and unattributed. The text-based methods of Section 1.4 recover latent positions rather than topic-specific stances over time. Finally, consensus has been simulated either as convergence (Section 1.5) or, where learning-based methods were used, on synthetic data rather than the institution’s own deliberative records (Section 1.6). These literatures differ in method but none recovers a directional preference: who is for or against what, extracted from the regime’s own words.

This gap has two parts. The knowledge gap is the question raised in Section 1.1: whether the regime’s outputs, including failures such as the 2022 ATCM’s inability to agree its final report, can be anticipated from the parties’ positions at all. The methodological gap is the missing instrument: no existing method measures each party’s directional stance, per topic and over time, from the regime’s own records. Without that instrument the question cannot be answered. Section 1.8 turns this gap into three research questions.

No prior work extracts a directional preference for each party over time from a consensus regime’s own textual records, nor uses such preferences to drive a strategic simulation

in which agents have veto power. This thesis supplies both halves and joins them: the preferences recovered from the record become the parameters of the simulation that reproduces consensus and its failures.

1.8 Research problem

The first goal of this thesis is to reconstruct parties' preferences from the written records into a directional measure (stance) across topics and time. The second goal is to extract a directional opposition (veto) signal from the same records, and to test whether the combined preferences recover known history and predict decisions. Finally, this thesis aim to use these preferences to parameterise a multi-agent model that simulates consensus and its failures.

Chapter 2

Program and design of the research project

2.1 Objectives, methodology and research plan

This thesis has three interconnected goals. The first is to reconstruct each party's preferences from the regime's written records, as a directional measure across topics and time (Chapter 1). The second is to test whether the recovered preferences are credible: they must match known historical positions and align with the ATCM decision record (Chapter 2). The third is to use the validated preferences to simulate the regime's collective decision-making, so that its outputs, including consensus failures, can be anticipated (Chapter 3). Figure 2.1 shows the overall flow of the research project, the output of each stage, and how the objectives relate to one another. Throughout the following sections each chapter is described in detail.

2.1.1 Objective 1 / Chapter 1: The Support layer

Aim To produce a defensible, directional measure of parties' preferences over relevant topics through time, using documents the parties submit themselves. For the ATCM case, Working Papers (WPs) are the input data because, of all submitted documents, these are the only ones submitted by parties or Observers (non-Secretariat) that require discussion and action. This makes WPs the most likely text to contain parties' strategic intentions. They are also where outcomes originate: nothing happens without them, and Measures generally begin as WPs. The analysis uses the English-language WPs available in the Secretariat's public database. Since WPs are translated into the four treaty languages (Secretariat of the Antarctic Treaty 2023), the English corpus is not limited to papers originally drafted in English. It is also practical: most of the available




Figure 2.1: Thesis workflow: from text records to a validated consensus simulation.

pre-trained embedding and NLI models were trained on English-language data.

Approach The approach is a pipeline with multiple stages. The first set of stages discovers what the corpus is about. Each document is embedded: mapped to a vector in $\mathbb{R}^m$, the embedding space introduced in Section 1.4. The assumption carried by this step is geometric: documents with similar meaning receive nearby vectors. BERTopic (Grootendorst 2022) then reduces the dimension of these vectors and clusters them, and each cluster is a topic. The output is the set of topics K with cluster centroids $\{\boldsymbol{\mu}_k\}_{k \in K}$, where each $\boldsymbol{\mu}_k \in \mathbb{R}^m$ is the average embedding of the documents in topic k . The second set of stages selects, for each document and topic, the passages relevant to that topic. Each document d is split into passages $P_d = \{p_1, \dots, p_{n_d}\}$, and each passage p is embedded as $\mathbf{e}_p \in \mathbb{R}^m$. Relevance is defined quantitatively: the relevance of passage p to topic k is the cosine similarity between the passage embedding and the topic centroid,

$$\rho_{p,k} = \cos(\mathbf{e}_p, \boldsymbol{\mu}_k) = \frac{\mathbf{e}_p \cdot \boldsymbol{\mu}_k}{\|\mathbf{e}_p\| \|\boldsymbol{\mu}_k\|}.$$

Under the geometric assumption above, a passage that points in the same direction as the topic centroid discusses that topic. Relevance then acts as a gate: $P_{d,k}$ is the set of passages scoring above a threshold, and only these pass to the next stage. Then the topic-relevant sentences in each $P_{d,k}$ are extracted using a Large Language Model (LLM); $\text{Extract}(P_{d,k})$ denotes this step. The extracted sentences are concatenated and scored as a set by the NLI classifier introduced in Section 1.4: given the text as premise and a hypothesis about topic k , it scores entailment, neutrality, and contradiction. Entailment means the premise gives support to the hypothesis. Contradiction means the premise contradicts it. The resulting scores are binned and combined into five ordinal classes, producing a distribution $\pi_{d,k}$ that ranges from strong opposition to strong support. Formally, it is as follows

$$\pi_{d,k} = \text{NLI}(\text{Extract}(P_{d,k})) \in \Delta_5, \quad \Delta_5 = \left\{ \pi \in \mathbb{R}^5 \mid \pi_\ell \geq 0, \sum_{\ell=1}^5 \pi_\ell = 1 \right\}, \quad (2.1)$$

Here Δ_5 is the probability simplex: $\pi_{d,k}$ assigns a non-negative probability to each of the five classes, and these probabilities sum to one. The signed stance of document d on topic k is the expected value of this distribution under an equal-spacing convention. This is:

$$s_{d,k} = \sum_{\ell=1}^5 v_\ell \pi_{d,k}(\ell), \quad v = \left(-1, -\frac{1}{2}, 0, \frac{1}{2}, 1 \right). \quad (2.2)$$

Since $\pi_{d,k}$ is a probability distribution and the five fixed class values v_ℓ all lie in $[-1, 1]$, the stance $s_{d,k}$ is a convex combination of those values. The stance is therefore a

continuous quantity in $[-1, 1]$ on topic k 's axis, while v itself never varies.

Finally, the previous result is aggregated to produce country $\times$ topic $\times$ year stances over the Working Papers, weighted by co-sponsorship credit. This aggregation is a design choice of this thesis rather than a result from prior work: the weights below are conventions, each chosen to answer a different question.

$$S_{c,k,y} = \sum_{\substack{d \in D(c,y) \\ k \in K(d)}} w_{d,c} s_{d,k}, \quad (2.3)$$

where $D(c, y)$ is the set of Working Papers co-authored by country c in year y , $K(d)$ is the set of topics assigned to document d , and $w_{d,c}$ is the attribution weight of that country in one document. Two attribution conventions are computed:

- Full credit, $w_{d,c} = 1$: every co-sponsor is counted as fully holding the position. This is the coalition-membership reading, and it measures which states align on a topic (“who is on this side?”).
- Fractional credit, $w_{d,c} = 1/|A_d|$, with A_d the set of co-sponsors of d : a paper’s positional mass is divided among its sponsors, so a state that authors alone carries more weight than one of eight co-sponsors. This is the influence reading, and it discounts states that merely sign on (“who is driving this?”).

For a fixed year y , the slice $S_{\cdot,\cdot,y} \in \mathbb{R}^{C \times |K|}$ is a signed country $\times$ topic matrix, with C countries and $|K|$ topics. A dot in place of an index means that the index runs over all its values, so the slice collects every country’s stance on every topic in year y . See Appendix C for full notation.

Measuring directional stance in diplomatic text is challenging because diplomats hedge. Hedging here means cautious, non-committal phrasing: a party rarely writes “we oppose this measure”, it writes “further consideration may be warranted”. Positions on opposite sides also share the same vocabulary of safeguards: a phrase such as “subject to appropriate environmental safeguards” appears in permissive and restrictive papers alike. If the classifier scores a single hypothesis, this hedged language pulls every score toward neutrality.

Stance is therefore recovered as a contrast. For each topic k , three hypotheses are written: one at the supporting pole (h_k^+ , e.g. “tourism should be further regulated”), one neutral (h_k^0), and one at the opposing pole (h_k^- , e.g. “tourism regulation should be relaxed”). Each hypothesis is paired with the extracted text as premise and scored by the NLI classifier (Laurer et al. 2024). Subtracting the against-pole score from the for-pole score gives a direction that survives the hedged phrasing. As a robustness

check, the same texts are scored a second time by a generative LLM used as a secondary instrument (Le Mens and Gallego 2025), and the two instruments are compared.

Validation strategy The pipeline above is an assertion built on stated assumptions: that meaning is geometric, that relevance is cosine similarity to a topic centroid, that stance is a contrast of entailment scores, and that co-sponsors share a paper’s position. None of these is guaranteed to hold on ATCM text. The construct is therefore validated with three independent checks, each targeting a different stage of the pipeline. A topic validation check compares the discovered topics against the Secretariat’s categories. Every WP carries a category assigned by the Secretariat, so each paper has both an official category and a discovered topic. The two labellings are cross-tabulated, and their agreement is measured by majority-category overlap: the share of papers in each discovered topic that belong to that topic’s most common official category. A historical canary uses the minerals/CRAMRA negotiations of the late 1970s and 1980s as documented ground truth. The parties’ positions in that episode are recorded in the historical literature (Jackson 2021): most prominently, Australia and France abandoned the negotiated convention and pushed instead for a mining ban. The test checks whether the pipeline recovers those documented directions from the papers alone. Finally, a co-sponsorship check tests the assumption that countries co-sponsor a WP because they share its topic and stance: within the same year and topic, countries whose recovered stances are similar should be the ones that actually co-sponsor papers, with stance similarity computed excluding the paper being predicted.

Figure 2.2 below gives a summarised overview of the steps described above.

2.1.2 Objective 2 / Chapter 2: The Opposition layer and external validation

Aim To measure not only what parties support but what they would block. Consensus fails through objection, so anticipating failure requires knowing who opposes what. This is the opposition signal: a directional measure of a party’s resistance to a proposal. It is not the inverse of the support layer, because the two capture different behaviours. The support layer already contains disagreement of one kind: a party proposes one direction while another proposes the contrary. The veto signal is different: it is the act of objecting to a proposal on the table. A party may hold the opposite stance on a topic yet never block, and blocking is rarer and more consequential. The signal has two sides. The first is opposition to a given proposal itself: the *veto signal*. The second is a secondary, exploratory result: a country’s general propensity to reject proposals within a given ATCM.




Figure 2.2: Chapter 1 support-layer pipeline, from text records to the country×topic×year support tensor S .

Approach As with the support layer, this chapter builds a pipeline and, for the specific case of the ATS, the ATCMs’ final reports are used as text input. Given the scarcity of text in which countries veto a proposal, the attribution is expected to be harder than for the support layer. Each extracted statement yields a *(country, target, polarity)* triple, with rationale and source quote. Taking the topics K from Chapter 1 and mapping targets back to them yields an opposition tensor $\mathbf{O}$ with the same Country $\times$ Topic $\times$ Year structure. At this point, there is a pair $(\mathbf{S}, \mathbf{O})$ of tensors that constitutes a signed preference network. These two tensors are the inputs to the third chapter.

External validation The combined structure (the tensor pair $(\mathbf{S}, \mathbf{O})$) is tested for predictive content against observed outcomes (see Section 1.2) and the compilation by Gardiner et al. (2025), for the case of the ATS. This gives Chapter 2 an independent research question, tests the results so far, and produces a quantitative description. The novelty in this work is the directional signed signal derived from the institution’s own input records and outcomes. Although this thesis focuses on the ATS, the pipeline is institution-agnostic: the same $(\mathbf{S}, \mathbf{O})$ construction could be applied to other consensus- or veto-based international institutions whose deliberative records are available.

The main risks in this chapter are the sparsity of text that can be directly attributed to an input, and the identification of the country that vetoes any proposal. A further limitation is the influence of external geopolitics, as in the case of the consensus breakdown in 2022 due to the war in Ukraine (Gardiner et al. 2025). The cause of this type of exogenous influence is unlikely to be reflected in the text data of an institution not directly related to such issues.

Figure 2.3 sketches the opposition layer pipeline and its external-validation step for the ATCM case study.

2.1.3 Objective 3 / Chapter 3: The Simulation layer

Aim To simulate consensus processes with country-agents whose state is parameterised by the empirically derived support and opposition from previous chapters. Its main measure of success is to be able to reproduce historical trajectories, and examine institutional counterfactuals. In consensus-based systems every Party is a veto player (Tsebelis 2002) and the representation is a strategic bargaining game over a status quo rather than an averaging dynamic.

Approach To tackle this task, a custom negotiation environment is specified in which proposals are presented one at a time. Each proposal is an input text/Working Paper d




Figure 2.3: Chapter 2 opposition layer sketch: from text data to opposition tensor O , forming the signed preference network (S, O) .

that is mostly about topic k in year y . Then the agents (one per consultative party) decide whether or not to object to a proposal. Each agent's decision-making is supported by its support vector (Chapter 1), its opposition history (Chapter 2), and metadata such as consultative status, official language, and GDP, used as context features rather than as an independent contribution. An agent's action on a presented proposal is binary: to object or not, $o_i \in \{0, 1\}$. Under the consensus rule the proposal is adopted only if no Party objects,

$$1_{\text{adopt}} = \prod_j (1 - o_j), \quad (2.4)$$

therefore, a single objection vetoes adoption. The reward to agent i is then

$$R_i = \alpha S_{i,k,y} 1_{\text{adopt}} - \beta O_{i,k,y} 1_{\text{adopt}} - \gamma o_i, \quad (2.5)$$

where the first term rewards adoption of outcomes the Party supports. The second penalises adoption of outcomes it opposes. The third subtracts a diplomatic cost for casting an objection. Finally, $\alpha, \beta, \gamma \geq 0$ are non-negative trade-off weights. R_i is measured relative to the status quo and normalised to zero. For a non-zero 1_{adopt} all parties need to agree (non-object), therefore, every agent holds a genuine veto. Furthermore, consensus emerges from the strategic trade-off between blocking a proposal and bearing the cost of objecting.

The model is validated by the predicted consensus or failure and by an ablation that removes different layers to quantify its contribution. Then a counterfactual analysis examines alternative decision rules, for example, qualified majority in place of unanimity, which relaxes (2.4). This counterfactual intends to predict the behaviour of the system given a change in the institutional rules e.g. ATCM framework (Secretariat of the Antarctic Treaty 2023). This is compared against the effectiveness metrics of Gardiner et al. (2025). This chapter trains its agents with empirically extracted positions from the institution's own record. This methodology makes this the highest-risk chapter and is framed as a proof of concept, contingent on Chapters 1 and 2 passing validation.

2.2 Resources and funding required

The project uses openly published data from the Antarctic Treaty Secretariat document and may use other public data by other international institutions. The project uses open-source software and/or QUT-licensed software. Computation is served by the QUT high-performance computing cluster. No additional funding is required.

2.3 Individual contribution to the research team

I am the sole researcher on this project, with guidance from my supervisory team. The Chapter 1 infrastructure developed at the outset of candidature is preliminary work that establishes feasibility. This is not presented in this document as thesis results, and all reported findings will be produced within candidature. The thesis extends prior group work by Botelho et al. (2026), Kumar et al. (2026) and Carter et al. (2026).

2.4 Acknowledging use of Artificial Intelligence (AI) and AI-assisted technologies

I have used Claude Code for boilerplate and scaffolding code in the implementation of the pipelines and scripts. All architectural decisions, methodological choices, and algorithm design were specified by me in the input prompts. I also reviewed and edited generated code. I have also used Consensus AI and the Claude app to fetch the most relevant scientific papers from the available literature, based on my own ideas, keywords, and criteria. Afterwards, I read the papers and drew conclusions myself. Finally, I used Claude Code as a writing assistant for this document, to help with grammar, readability, clarity and LaTeX document formatting. All research questions, project design, ideas, interpretation of results, and selection of literature are my own.

2.5 Timeline for completion

The candidature is full-time and commenced on the 31st of March, 2026. All milestones are placed against the QUT full-time schedule (Table 2.1).

Table 2.1: Candidature timeline (April 2026–April 2029). ■ Ch. 1 ■ Ch. 2 ■ Ch. 3 ■ Milestones/events.



Table 2.1 (continued): Conference Attendance and Coursework.

Activity	3	6	9	12	15	18	21	24	27	30	33	36
	Jul	Oct	Jan	Apr	Jul	Oct	Jan	Apr	Jul	Oct	Jan	Apr
	'26	'26	'27	'27	'27	'27	'28	'28	'28	'28	'29	'29
Conference Attendance												
QANZIAM 2026												
AMEG Symposium 2026												
QANZIAM 2027												
International Conference												
AMEG Symposium 2027												
QANZIAM 2028												
Coursework												
RIO (completed)												
AIRS / IFN006												

Coursework (IFN006 Advanced Information Research Skills) is currently underway and will be completed by July 2026 following Stage 2 submission. The Research Integrity Online quiz is completed (see Appendix A).

References

- Abercrombie, Gavin and Riza Batista-Navarro (2022). “Policy-focused stance detection in parliamentary debate speeches”. In: *Northern European Journal of Language Technology* 8.1. DOI: [10.3384/nejlt.2000-1533.2022.3264](https://doi.org/10.3384/nejlt.2000-1533.2022.3264).
- Bachrach, Yoram, Richard Everett, Edward Hughes, Angeliki Lazaridou, Marc Lanctot, Guy Lever, Michael Vlassopoulos, and Thore Graepel (2020). “Negotiating team formation using deep reinforcement learning”. In: *Artificial Intelligence* 288, p. 103356. DOI: [10.1016/j.artint.2020.103356](https://doi.org/10.1016/j.artint.2020.103356).
- Baturo, Alexander, Niheer Dasandi, and Slava J. Mikhaylov (2017). “Understanding state preferences with text as data: Introducing the UN General Debate corpus”. In: *Research & Politics* 4.2. DOI: [10.1177/2053168017712821](https://doi.org/10.1177/2053168017712821).
- Benac, Leo, Jonas Raedler, Zilin Ma, and Finale Doshi-Velez (2026). *A Benchmark for Multi-Party Negotiation Games from Real Negotiation Data*. DOI: [10.48550/arXiv.2603.14066](https://doi.org/10.48550/arXiv.2603.14066). arXiv: [2603.14066](https://arxiv.org/abs/2603.14066).
- Bernardo, Carmela, Lingfei Wang, Francesco Vasca, Yiguang Hong, Guodong Shi, and Claudio Altafini (2021). “Achieving consensus in multilateral international negotiations: The case study of the 2015 Paris Agreement on climate change”. In: *Science Advances* 7.51, eabg8068. DOI: [10.1126/sciadv.abg8068](https://doi.org/10.1126/sciadv.abg8068).
- Bolleddu, Divya (2025). *Dialogue Diplomats: An end-to-end multi-agent reinforcement learning system for automated conflict resolution and consensus building*. DOI: [10.48550/arXiv.2511.17654](https://doi.org/10.48550/arXiv.2511.17654). arXiv: [2511.17654](https://arxiv.org/abs/2511.17654).
- Botelho, Larissa Lubiana, Kate J. Helmstedt, Kerrie A. Wilson, and Michael Bode (2026). “Coalition Dynamics in Antarctic Governance”. Unpublished manuscript.
- Burnham, Michael (2023). “Stance detection: A practical guide to classifying political beliefs in text”. In: *Political Science Research and Methods*. DOI: [10.1017/psrm.2024.35](https://doi.org/10.1017/psrm.2024.35).
- Carter, Zachary T., Larissa Lubiana Botelho, Steven L. Chown, Parsa Yaryab, Nadiah P. Kristensen, Kerrie A. Wilson, Daniela Liggett, and Michael Bode (2026). “The Limits of Antarctic Conservation Governance”. Manuscript submitted for publication.

- Chang, Ho-Chun Herbert (2021). “Multi-issue negotiation with deep reinforcement learning”. In: *Knowledge-Based Systems* 211, p. 106544. DOI: [10.1016/j.knosys.2020.106544](https://doi.org/10.1016/j.knosys.2020.106544).
- DeGroot, Morris H. (1974). “Reaching a consensus”. In: *Journal of the American Statistical Association* 69.345, pp. 118–121. DOI: [10.1080/01621459.1974.10480137](https://doi.org/10.1080/01621459.1974.10480137).
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2019). “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”. In: *Proceedings of NAACL-HLT 2019*, pp. 4171–4186. DOI: [10.18653/v1/N19-1423](https://doi.org/10.18653/v1/N19-1423).
- Friedkin, Noah E. and Eugene C. Johnsen (1990). “Social Influence and Opinions”. In: *Journal of Mathematical Sociology* 15.3–4, pp. 193–206. DOI: [10.1080/0022250X.1990.9990069](https://doi.org/10.1080/0022250X.1990.9990069).
- Gardiner, Nicholas B., Natasha Gilbert, Daniela Liggett, and Michael Bode (2025). “Measuring the performance of Antarctic Treaty decision-making”. In: *Conservation Biology* 39, e14349. DOI: [10.1111/cobi.14349](https://doi.org/10.1111/cobi.14349).
- Gray, Julia (2018). “Life, Death, or Zombie? The Vitality of International Organizations”. In: *International Studies Quarterly* 62.1, pp. 1–13. DOI: [10.1093/isq/sqx086](https://doi.org/10.1093/isq/sqx086).
- Grimmer, Justin and Brandon M. Stewart (2013). “Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts”. In: *Political Analysis* 21.3, pp. 267–297. DOI: [10.1093/pan/mps028](https://doi.org/10.1093/pan/mps028).
- Gronauer, Sven and Klaus Diepold (2022). “Multi-agent deep reinforcement learning: A survey”. In: *Artificial Intelligence Review* 55, pp. 895–943. DOI: [10.1007/s10462-021-09996-w](https://doi.org/10.1007/s10462-021-09996-w).
- Grootendorst, Maarten (2022). *BERTopic: Neural Topic Modeling with a Class-Based TF-IDF Procedure*. arXiv: [2203.05794](https://arxiv.org/abs/2203.05794) [cs.CL]. URL: <https://arxiv.org/abs/2203.05794>.
- Hegselmann, Rainer and Ulrich Krause (2002). “Opinion Dynamics and Bounded Confidence: Models, Analysis and Simulation”. In: *Journal of Artificial Societies and Social Simulation* 5.3. URL: <https://www.jasss.org/5/3/2.html>.
- Jackson, Andrew (2021). *Who Saved Antarctica?* Cham: Palgrave Macmillan. ISBN: 978-3-030-78404-1. DOI: [10.1007/978-3-030-78405-8](https://doi.org/10.1007/978-3-030-78405-8).
- Kulkarni, Abhishek N., Andy Liu, Jean-Raphaël Gaglione, Daniel Fried, and Ufuk Topcu (2025). “Dynamic Coalition Structure Detection in Natural Language-based Interactions”. In: *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. Detroit, MI, USA: IFAAMAS, pp. 1227–1234.
- Kumar, Viney, Valeria Senigaglia, Maria Kleshnina, and Michael Bode (2026). “Consistent Patterns of Interaction Within Antarctic Governance Forums”. Unpublished manuscript.

- Lauderdale, Benjamin E. and Alexander Herzog (2016). “Measuring political positions from legislative speech”. In: *Political Analysis* 24.3, pp. 374–394. DOI: [10.1093/pan/mpw017](https://doi.org/10.1093/pan/mpw017).
- Laurer, Moritz, Wouter van Atteveldt, Andreu Casas, and Kasper Welbers (2024). *Building Efficient Universal Classifiers with Natural Language Inference*. DOI: [10.48550/arXiv.2312.17543](https://doi.org/10.48550/arXiv.2312.17543). arXiv: [2312.17543](https://arxiv.org/abs/2312.17543) [cs.CL]. URL: <https://arxiv.org/abs/2312.17543>.
- Le Mens, Gaël and Aina Gallego (2025). “Positioning Political Texts with Large Language Models by Asking and Averaging”. In: *Political Analysis* 33.3, pp. 274–282. DOI: [10.1017/pan.2024.29](https://doi.org/10.1017/pan.2024.29).
- Li, Zun, Marc Lanctot, Kevin R. McKee, Luke Marris, et al. (2023). *Search-improved game-theoretic multiagent reinforcement learning in general and negotiation games*. arXiv preprint. arXiv: [2305.11607](https://arxiv.org/abs/2305.11607).
- Littman, Michael L. (1994). “Markov Games as a Framework for Multi-agent Reinforcement Learning”. In: *Proceedings of the Eleventh International Conference on Machine Learning (ICML)*, pp. 157–163.
- Mancilla, Alejandra and Julia Ann Jabour (2022). “Turned 60, is the Antarctic Treaty System in good health?” In: *The Geographical Journal* 188.CONFIRM, CONFIRM. DOI: [10.1111/geoj.12501](https://doi.org/10.1111/geoj.12501).
- Moghimifar, Farhad, Yuan-Fang Li, Robert Thomson, and Gholamreza Haffari (2024). *Modelling political coalition negotiations using LLM-based agents*. DOI: [10.48550/arXiv.2402.11712](https://doi.org/10.48550/arXiv.2402.11712). arXiv: [2402.11712](https://arxiv.org/abs/2402.11712).
- Parschan, Patrick and Charlott Frederike Jakob (2025). “Computational measurement of political positions: A review of text-based ideal point estimation algorithms”. In: *Quality & Quantity*. DOI: [10.1007/s11135-025-02500-4](https://doi.org/10.1007/s11135-025-02500-4). arXiv: [2511.13238](https://arxiv.org/abs/2511.13238).
- Pérez, I. J., F. J. Cabrerizo, S. Alonso, Yucheng Dong, F. Chiclana, and E. Herrera-Viedma (2018). “On Dynamic Consensus Processes in Group Decision Making Problems”. In: *Information Sciences* 459, pp. 165–181. DOI: [10.1016/j.ins.2018.05.017](https://doi.org/10.1016/j.ins.2018.05.017).
- Rheault, Ludovic and Christopher Cochrane (2019). “Word embeddings for the analysis of ideological placement in parliamentary corpora”. In: *Political Analysis* 28.1, pp. 112–133. DOI: [10.1017/pan.2019.26](https://doi.org/10.1017/pan.2019.26).
- Rietig, Katharina, Christina Peringer, Sarina Theys, and Jared Censoro (2023). “Unanimity or Standing Aside? Reinterpreting Consensus in United Nations Framework Convention on Climate Change Negotiations”. In: *International Environmental Agreements: Politics, Law and Economics*. DOI: [10.1007/s10784-023-09593-y](https://doi.org/10.1007/s10784-023-09593-y).
- Rodriguez, Pedro L. and Arthur Spirling (2022). “Word embeddings: What works, what doesn’t, and how to tell the difference for applied research”. In: *The Journal of Politics* 84.1, pp. 101–115. DOI: [10.1086/715162](https://doi.org/10.1086/715162).

- Sarkar, Samriddhi, Mariana Curado Malta, and Animesh Dutta (2022). “A survey on applications of coalition formation in multi-agent systems”. In: *Concurrency and Computation: Practice and Experience* 34.11, e6876. DOI: <https://doi.org/10.1002/cpe.6876>.
- Schmid, Kyrill, Lenz Belzner, Thomy Phan, Thomas Gabor, and Claudia Linnhoff-Popien (2020). “Multi-agent Reinforcement Learning for Bargaining under Risk and Asymmetric Information”. In: pp. 144–151. DOI: [10.5220/0008913901440151](https://doi.org/10.5220/0008913901440151).
- Secretariat of the Antarctic Treaty (2023). *Rules of Procedure of the Antarctic Treaty Consultative Meeting and the Committee for Environmental Protection*. Updated August 2023. Buenos Aires: Secretariat of the Antarctic Treaty. 46 pp. ISBN: 978-987-8929-23-1. URL: <https://www.ats.aq> (visited on 06/09/2026).
- Slapin, Jonathan B. and Sven-Oliver Proksch (2008). “A scaling model for estimating time-series party positions from texts”. In: *American Journal of Political Science* 52.3, pp. 705–722. DOI: [10.1111/j.1540-5907.2008.00338.x](https://doi.org/10.1111/j.1540-5907.2008.00338.x).
- Sykora-Bodie, Seth T. and Tiffany H. Morrison (2019). “Drivers of Consensus-Based Decision-Making in International Environmental Regimes: Lessons from the Southern Ocean”. In: *Aquatic Conservation: Marine and Freshwater Ecosystems* 29.12, pp. 2147–2161. DOI: <https://doi.org/10.1002/aqc.3200>.
- Tran, Thi Ngoc Trang, Alexander Felfernig, and Viet-Man Le (2024). “An Overview of Consensus Models for Group Decision-Making and Group Recommender Systems”. In: *User Modeling and User-Adapted Interaction* 34, pp. 489–547. DOI: [10.1007/s11257-023-09380-z](https://doi.org/10.1007/s11257-023-09380-z).
- Tsebelis, George (2002). *Veto Players: How Political Institutions Work*. Princeton, NJ: Princeton University Press.
- Vafa, Keyon, Suresh Naidu, and David M. Blei (2020). “Text-based ideal points”. In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 5345–5357. DOI: [10.18653/v1/2020.acl-main.475](https://doi.org/10.18653/v1/2020.acl-main.475).

Appendices

Appendix A: Research Integrity Online

[Insert RIO completion certificate. I am on it, next I will add it to this document before submission]

Appendix B: Coursework

AIRS / IFN006 is currently underway, this course will end on july, 2026.

Appendix C: Notation

	c, i	countries (Consultative Parties), $c \in \{1, \dots, C\}$
	k	topic, $k \in K$; $ K $ is the number of topics
	y	year
	d	document (Working Paper)
	p	passage (chunk) of a document
Indices and sets.	ℓ	ordinal stance class, $\ell \in \{1, \dots, 5\}$
	$D(c, y)$	Working Papers co-sponsored by country c in year y
	$K(d)$	topics assigned to document d
	P_d	passages of document d , $P_d = \{p_1, \dots, p_{n_d}\}$
	$P_{d,k}$	passages of d selected as relevant to topic k
	A_d	co-sponsors of document d
Embedding space.	m	embedding dimensionality (full Stella space)
	$\mathbf{e}_p$	embedding of passage p , $\mathbf{e}_p \in \mathbb{R}^m$
	$\boldsymbol{\mu}_k$	centroid of topic k , $\boldsymbol{\mu}_k \in \mathbb{R}^m$
	$\rho_{p,k}$	passage–topic relevance, $\rho_{p,k} = \cos(\mathbf{e}_p, \boldsymbol{\mu}_k)$
	τ	relevance selection threshold

	h_k^+, h_k^0, h_k^-	for-, neutral-, against-pole hypotheses for topic k
	$\lambda(d, h)$	null-premise-debiased entailment log-ratio (DCPMI)
	$\tilde{\pi}_{d,k}$	three-pole distribution (stance classifier)
Stance and hypotheses.	$\pi_{d,k}$	five-class ordinal distribution, $\pi_{d,k} \in \Delta_5$
	Δ_5	probability simplex, $\{\pi \in \mathbb{R}^5 : \pi_\ell \geq 0, \sum_\ell \pi_\ell = 1\}$
	v	class values, $v = (-1, -\frac{1}{2}, 0, \frac{1}{2}, 1)$
	$s_{d,k}$	signed stance of d on topic k , $s_{d,k} \in [-1, 1]$
Aggregation and tensors.	$w_{d,c}$	attribution weight: 1 (full credit) or $1/ A_d $ (fractional)
	$S_{c,k,y}$	support of country c on topic k in year y , Eq. (2.3)
	$S_{\cdot,\cdot,y}$	year- y stance slice, $S_{\cdot,\cdot,y} \in \mathbb{R}^{C \times K }$
	$O_{c,k,y}$	opposition of country c on topic k in year y

For these tensors, high sparsity is expected because absence of evidence is recorded as missing, not as zero.

Stance classifier (Chapter 1). For topic k , let h_k^+, h_k^0, h_k^- be the for-, neutral-, and against-pole hypotheses. For passage set d and hypothesis h , the debiased evidence is the null-premise-corrected log-ratio,

$$\lambda(d, h) = \text{clip}\left(\log P(\text{entail} \mid d, h) - \log P(\text{entail} \mid \emptyset, h), \pm 5\right),$$

this removes each hypothesis's marginal acceptability. The three-pole distribution combines entailment of each pole with contradiction of its opposite,

$$\tilde{\pi}_{d,k} = \text{softmax}\begin{pmatrix} \lambda_{\text{ent}}(h_k^+) + \lambda_{\text{con}}(h_k^-) \\ \lambda_{\text{ent}}(h_k^0) \\ \lambda_{\text{ent}}(h_k^-) + \lambda_{\text{con}}(h_k^+) \end{pmatrix},$$

these values are expanded to the five ordinal classes $\pi_{d,k}$ and reduced to the scalar stance by the expected value of Equation 2.2.